

# ChandonnetPythonProject

December 16, 2022

## 1 Final Python Project

### 1.0.1 Ray Chandonnet

### 1.0.2 12/16/2022

Key Assumptions and modeling approaches:

Assumptions: - All salaries in the Salary file are already in USD and so no FX conversion was done (this wasn't totally clear but seemed the case) - Only used salary for jobs titled as Data Scientist - Assigned experience levels as follows: ~ 0-5 years experience = "Junior" ~ >5 years experience = "Senior" - Used all salaries for all dates, meaning I assume the data is pretty evenly distributed across years for all cities (Otherwise data is actually not granular enough to be meaningful) - I only use CASH compensation to calculate affordability, since you can't pay bills with stock grants, which are intended to build wealth and usually are on a vesting schedule

Modeling approaches:

- I consider the most important index to be "Cost of Living Plus Rent" which factors in housing costs which are a critical component of cost of living
- The cost of living indexes provided use NYC as the baseline index of 100, so all other indexes are expressed relative to NYC
- Given that, I construct a comparable "salary index" for all cities in the salary file, also relative to NYC. So NYC is 100; if a job pays only 75% less than the average comparable job in NYC it gets a salary index of 75; If it pays 1.2x a comparable job in NYC it gets a salary index of 120
- These indexes are done using separate average salaries for the two experience levels I created, to avoid skewing results where more junior or senior people exist in a given city, making salaries seem higher or lower than they really are
- I then create an "affordability index" which is salary index divided by cost of living index. Since both indices are relative to NYC, NYC is deemed to have an affordability index of 100, with more affordable cities than NYC having an index higher than 100 and less affordable cities than NYC having an index the result
- The idea is that how affordable a city is a function of how much you earn relative to the cost of living there. If an area costs a lot less to live than NYC, but pays almost as much, it's more affordable. The converse is also true.
- I then rank order the cities by affordability, for both junior and senior people.
- I also ran a few plots to see if there was a difference (skew) in overall affordability for junior vs senior people, and also ran

```
[27]: # import needed libraries and set file path prefix
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
path = "/users/raychandonnet/Dropbox (Personal)/Merrimack College - \
MS in Data Science/DSE5002/Python Project/"
```

## 2 Part 1: Data Wrangling the Cost of Living Index File

```
[28]: cost_of_living=pd.read_csv(path + "cost_of_living.csv")
cost_of_living=cost_of_living.rename(columns=\
    {"Cost of Living Plus Rent Index":"RealCOL_Index"})
# Step 1 - Break the City up - requires a lot of cleaning because city and
# country are in the same field separated by
# a comma, and US cities also have the state in form city, st, country.
# I want to break these up into separate columns to link to the country
# code data, then pull in job info for those city / country combos
# First, split the city into three separate fields
cost_of_living[['City', 'State', 'Country']]=cost_of_living['City'].str.\
    split(pat=",",n=2,expand=True)
# This leaves the city state and country correct for US cities, but has the
# country in the city column for non-US cities and "None" for country, so we
# have to fix that It's two steps:
# First, repalce all the "None" in Country column with the State value which
# is where the Country is. Then eliminate the dups by putting "-" in the State
# field for all those. So we basically just swapped values to get the country
# in the right place
cost_of_living['Country']=cost_of_living['Country'].\
    fillna(cost_of_living['State']) # Puts the state value in country where NA
cost_of_living['State']=np.where(cost_of_living['State']\
    ==cost_of_living['Country'],
    "-",cost_of_living['State'])
# It took me more time than I care to admit to realize that my states and
# countries had padded whitespace and so weren't matching when I tried to
# merge in the country code data!!! Sigh...using the strip method fixes that
cost_of_living['Country']=cost_of_living['Country'].str.strip()
cost_of_living['State']=cost_of_living['State'].str.strip()
# This leaves us with City, State and Country, with country spelled out
# Now I can merge in the country codes from that file into my table. This
# will let me connect to data in the salary and jobs files once I do some
# stuff with that data, if I need to use a code instead of the country
# spelled out
countrycodes=pd.read_excel(path+"country_codes.xlsx")
cost_of_living=pd.merge(left=cost_of_living,right=countrycodes,
    how='left',on='Country')
```

```
print(cost_of_living.head(10))
```

	Rank	City	Cost of Living Index	Rent Index	RealCOL_Index \
0	NaN	Hamilton	149.02	96.10	124.22
1	NaN	Zurich	131.24	69.26	102.19
2	NaN	Basel	130.93	49.38	92.70
3	NaN	Zug	128.13	72.12	101.87
4	NaN	Lugano	123.99	44.99	86.96
5	NaN	Lausanne	122.03	59.55	92.74
6	NaN	Beirut	120.47	27.76	77.01
7	NaN	Bern	118.16	46.12	84.39
8	NaN	Geneva	114.05	75.05	95.77
9	NaN	Stavanger	104.61	35.38	72.16

	Groceries Index	Restaurant Price Index	Local Purchasing Power Index \
0	157.89	155.22	79.43
1	136.14	132.52	129.79
2	137.07	130.95	111.53
3	132.61	130.93	143.40
4	129.17	119.80	111.96
5	122.56	127.01	127.01
6	141.33	116.95	15.40
7	118.37	120.88	112.46
8	112.70	126.31	120.60
9	102.46	107.51	85.90

	State	Country	Alpha-2 code	Alpha-3 code	Numeric
0	-	Bermuda	BM	BMU	60.0
1	-	Switzerland	CH	CHE	756.0
2	-	Switzerland	CH	CHE	756.0
3	-	Switzerland	CH	CHE	756.0
4	-	Switzerland	CH	CHE	756.0
5	-	Switzerland	CH	CHE	756.0
6	-	Lebanon	LB	LBN	422.0
7	-	Switzerland	CH	CHE	756.0
8	-	Switzerland	CH	CHE	756.0
9	-	Norway	NO	NOR	578.0

### 3 Part 2 - Data wrangling the salary data

I chose to use this data as the source for my income in calculating salary indexes relative to NY, a necessary pre-requisite to calculating an affordability index. There so many different ways we could slice and dice the data - I will be focusing on cash compensation and seniority, specifically for the salaries listed there as Data Scientist

```
[29]: # Read in the salaries file and do some data wrangling here as well to
      # convert the time stamp to a date, and put the cities and countries in the
```

```

# right columns
#
# I did not end up using the date column in any way - ran out of time
# But wanted to show I knew how to convert it
salaries=pd.read_csv(path+"Levels_Fyi_Salary_Data.csv") # read data
salaries=salaries[salaries['title']=="Data Scientist"] # Slice the DS jobs
salaries['timestamp'] = pd.to_datetime(salaries['timestamp']) # convert date
salaries=salaries.rename({"timestamp":"date"},axis=1) #rename date column
salaries[['City','State','Country']]=salaries['location'].\\
    str.split(pat=" ",n=2,expand=True)
# Now fill in "United States" where country is blank, delete the state for
# non-US locations, and remove whitespace
salaries['Country']=salaries['Country'].fillna("United States")
salaries['State']=np.where(salaries['Country']!='United States',
    "-",salaries['State']) # delete states for non-US
salaries['Country']=salaries['Country'].str.strip()
salaries['State']=salaries['State'].str.strip()
print(salaries.head(10))

```

	date	company	level	title \
419	2018-06-05 14:06:30	LinkedIn	Senior	Data Scientist
440	2018-06-08 09:49:25	Microsoft	64	Data Scientist
444	2018-06-08 17:55:09	ebay	26	Data Scientist
454	2018-06-10 19:39:35	Twitter	Staff	Data Scientist
495	2018-06-17 11:39:38	Facebook	5	Data Scientist
499	2018-06-17 19:02:50	Amazon	L5	Data Scientist
509	2018-06-20 00:47:43	Microsoft	65	Data Scientist
510	2018-06-20 00:49:11	Google	L6	Data Scientist
513	2018-06-21 10:54:35	Netflix	Senior	Data Scientist
523	2018-06-25 08:45:29	Tesla	Senior Engineer	Data Scientist

	totalyearlycompensation	location	yearsofexperience \
419	233000	San Francisco, CA	4.0
440	218000	Seattle, WA	11.0
444	180000	San Jose, CA	10.0
454	500000	San Francisco, CA	4.0
495	370000	Seattle, WA	8.0
499	200000	Seattle, WA	3.0
509	340000	Bellevue, WA	11.0
510	690000	Kirkland, WA	10.0
513	600000	Los Gatos, CA	3.0
523	168000	Palo Alto, CA	8.0

	yearsatcompany	tag	basesalary	...	Race_Asian \
419	0.0	Data Analysis	162000.0	...	0
440	11.0	ML / AI	165000.0	...	0
444	5.0	NaN	0.0	...	0
454	4.0	ML / AI	200000.0	...	0

495	3.0		NaN	190000.0	...	0
499	0.0		ML / AI	150000.0	...	0
509	11.0		ML / AI	200000.0	...	0
510	0.0		ML / AI	240000.0	...	0
513	1.0		ML / AI	600000.0	...	0
523	3.0	Mechanical Engineering		118000.0	...	0

	Race_White	Race_Two_Or_More	Race_Black	Race_Hispanic	Race	Education \
419	0	0	0	0	NaN	NaN
440	0	0	0	0	NaN	NaN
444	0	0	0	0	NaN	NaN
454	0	0	0	0	NaN	NaN
495	0	0	0	0	NaN	NaN
499	0	0	0	0	NaN	NaN
509	0	0	0	0	NaN	NaN
510	0	0	0	0	NaN	NaN
513	0	0	0	0	NaN	NaN
523	0	0	0	0	NaN	NaN

	City	State	Country
419	San Francisco	CA	United States
440	Seattle	WA	United States
444	San Jose	CA	United States
454	San Francisco	CA	United States
495	Seattle	WA	United States
499	Seattle	WA	United States
509	Bellevue	WA	United States
510	Kirkland	WA	United States
513	Los Gatos	CA	United States
523	Palo Alto	CA	United States

[10 rows x 32 columns]

## 4 Part 3: Enriching the salary data

Now I'm going to enrich the data a bit to make it more granular and make more sense: First, I calculate "cashcomp" as cash compensation = base + bonus. Then I tag each salary with a "Junior" or "Senior" experience level as noted above

```
[30]: salaries['cashcomp']=salaries['basesalary']+salaries['bonus']
salaries=salaries[salaries['cashcomp']!=0] # eliminate zeros
salaries['explevel']=np.where(salaries['yearsofexperience']<=5,
                             "Junior","Senior")
print(salaries.head(10))
```

	date	company	level	title \
419	2018-06-05 14:06:30	LinkedIn	Senior	Data Scientist
440	2018-06-08 09:49:25	Microsoft	64	Data Scientist

454	2018-06-10 19:39:35	Twitter	Staff	Data Scientist
495	2018-06-17 11:39:38	Facebook	5	Data Scientist
499	2018-06-17 19:02:50	Amazon	L5	Data Scientist
509	2018-06-20 00:47:43	Microsoft	65	Data Scientist
510	2018-06-20 00:49:11	Google	L6	Data Scientist
513	2018-06-21 10:54:35	Netflix	Senior	Data Scientist
523	2018-06-25 08:45:29	Tesla	Senior Engineer	Data Scientist
535	2018-06-26 21:37:46	GrubHub	II	Data Scientist

	totalyearlycompensation	location	yearsofexperience	\
419	233000	San Francisco, CA	4.0	
440	218000	Seattle, WA	11.0	
454	500000	San Francisco, CA	4.0	
495	370000	Seattle, WA	8.0	
499	200000	Seattle, WA	3.0	
509	340000	Bellevue, WA	11.0	
510	690000	Kirkland, WA	10.0	
513	600000	Los Gatos, CA	3.0	
523	168000	Palo Alto, CA	8.0	
535	187000	New York, NY	4.0	

	yearsatcompany	tag	basesalary	...	\
419	0.0	Data Analysis	162000.0	...	
440	11.0	ML / AI	165000.0	...	
454	4.0	ML / AI	200000.0	...	
495	3.0	NaN	190000.0	...	
499	0.0	ML / AI	150000.0	...	
509	11.0	ML / AI	200000.0	...	
510	0.0	ML / AI	240000.0	...	
513	1.0	ML / AI	600000.0	...	
523	3.0	Mechanical Engineering	118000.0	...	
535	1.0	ML / AI	150000.0	...	

	Race_Two_Or_More	Race_Black	Race_Hispanic	Race	Education	\
419	0	0	0	NaN	NaN	
440	0	0	0	NaN	NaN	
454	0	0	0	NaN	NaN	
495	0	0	0	NaN	NaN	
499	0	0	0	NaN	NaN	
509	0	0	0	NaN	NaN	
510	0	0	0	NaN	NaN	
513	0	0	0	NaN	NaN	
523	0	0	0	NaN	NaN	
535	0	0	0	NaN	NaN	

	City	State	Country	cashcomp	explevel
419	San Francisco	CA	United States	172000.0	Junior
440	Seattle	WA	United States	188000.0	Senior

454	San Francisco	CA	United States	220000.0	Junior
495	Seattle	WA	United States	230000.0	Senior
499	Seattle	WA	United States	231000.0	Junior
509	Bellevue	WA	United States	260000.0	Senior
510	Kirkland	WA	United States	312000.0	Senior
513	Los Gatos	CA	United States	600000.0	Junior
523	Palo Alto	CA	United States	118000.0	Senior
535	New York	NY	United States	160000.0	Junior

[10 rows x 34 columns]

## 5 Part 4 - Aggregating by city/level, isolating NYC

Now I aggregate by city, state, country and experience level, calculating mean salary for each subset. Then, since in the cost of living data, the indexes re calculated using NYC as baseline (100 index), I calculate a “salary index” for each city and experience combination, relative to that same experience level in NYC so everything is on a common denominator

```
[31]: sal_by_city=salaries.groupby(['Country','City','State','explevel'],
                                   as_index=False).agg(avg_salary=('cashcomp',
                                                                np.mean))
NYCsal=salaries[salaries['City']=="New York"] # Extract all the NYC jobs
NYC_by_level=NYCsal.groupby('explevel').agg(
    NYC_avg_salary=('cashcomp', np.mean)) # Calc NYC mean salaries by level
print(sal_by_ctry.head(5))
print(sal_by_city.head(5))
print(NYC_by_level.head())
```

	Country	explevel	avg_salary
0	Australia	Junior	89333.333333
1	Australia	Senior	213500.000000
2	Austria	Junior	17000.000000
3	Brazil	Senior	28000.000000
4	Canada	Junior	95297.297297

	Country	City	State	explevel	avg_salary
0	Australia	Canberra	-	Junior	93000.0
1	Australia	Melbourne	-	Junior	67000.0
2	Australia	Melbourne	-	Senior	142000.0
3	Australia	Sydney	-	Junior	108000.0
4	Australia	Sydney	-	Senior	285000.0

	NYC_avg_salary
explevel	
Junior	156672.131148
Senior	230960.784314

## 6 Part 5 - Calculate salary and affordability indexes

Here comes the magic! Now that I have average salaries for NYC by level, I can use those to create a salary index, with NYC as the baseline at 100, for each city/level combo I found. Once that cost of living index is calculated, I can derive an “affordability” index dividing the salary index for that city/country/level into the cost of living index for that city/country. So for example if the salary index is 80 (80% of NYC) but the cost of living index is 40 (50% of NYC), then that location for that experience level is twice as affordable as NYC ( $80/40=2x$ ) or affordability = 200 meaning your money goes twice as far

```
[32]: afford_city_level=pd.merge(left=sal_by_city,right=NYC_by_level,
                                how='left',on='explevel') #Pull in NYC salary/level
afford_city_level['sal_index']=afford_city_level['avg_salary']\
    /afford_city_level['NYC_avg_salary']*100 # Calculate salary index
afford_city_level=pd.merge(left=afford_city_level,
                            right=cost_of_living,
                            how='left',
                            on=['City','State','Country'])
afford_city_level['affordability']=afford_city_level['sal_index']\
    /afford_city_level['RealCOL_Index']*100
print(afford_city_level.head())
```

	Country	City	State	explevel	avg_salary	NYC_avg_salary \
0	Australia	Canberra	-	Junior	93000.0	156672.131148
1	Australia	Melbourne	-	Junior	67000.0	156672.131148
2	Australia	Melbourne	-	Senior	142000.0	230960.784314
3	Australia	Sydney	-	Junior	108000.0	156672.131148
4	Australia	Sydney	-	Senior	285000.0	230960.784314

	sal_index	Rank	Cost of Living Index	Rent Index	RealCOL_Index \
0	59.359632	NaN	75.94	42.50	60.26
1	42.764466	NaN	76.76	38.65	58.90
2	61.482299	NaN	76.76	38.65	58.90
3	68.933766	NaN	83.21	58.03	71.41
4	123.397572	NaN	83.21	58.03	71.41

	Groceries Index	Restaurant Price Index	Local Purchasing Power Index \
0	76.81	79.07	105.11
1	77.78	74.68	102.16
2	77.78	74.68	102.16
3	79.65	73.06	104.52
4	79.65	73.06	104.52

	Alpha-2 code	Alpha-3 code	Numeric	affordability
0	AU	AUS	36.0	98.505861
1	AU	AUS	36.0	72.605205
2	AU	AUS	36.0	104.384209
3	AU	AUS	36.0	96.532371



## 7 Part 6 - Slice and analyze the results

I want to rank order affordability individually for junior and senior people, as well as overall. For overall, I will average the affordability index for junior and senior people. But I think it's more interesting to look at junior and senior people separately to see if there are big differences

```
[33]: # First delete all the NaN's, which represent cities for which we have no
# cost of living index and thus no affordability calculation
afford_city_level=afford_city_level[afford_city_level['affordability'].notna()]
# Now slice the data into Junior vs Senior people, and also average the
# affordability for each city/state/country
afford_city_junior=afford_city_level[afford_city_level['explevel']=="Junior"]
afford_city_junior=afford_city_junior[['Country', 'City', 'State',
                                         'RealCOL_Index', 'sal_index',
                                         'affordability']] # slice columns needed
afford_city_senior=afford_city_level[afford_city_level['explevel']=="Senior"]
afford_city_senior=afford_city_senior[['Country', 'City', 'State',
                                         'RealCOL_Index', 'sal_index',
                                         'affordability']] # slice columns needed
# Now group everything, averaging out the indexes for junior and senior people
afford_city_both=afford_city_level.groupby(['Country', 'City', 'State'],\
      as_index=False).agg(
      RealCOL_Index=('RealCOL_Index', np.mean),
      sal_index=('sal_index', np.mean),
      affordability=('affordability', np.mean))
# Sort all 3 tables descending by affordability
afford_city_junior=afford_city_junior.sort_values(by=['affordability'],
      ascending=False)
afford_city_junior=afford_city_junior.sort_values(by=['affordability'],
      ascending=False)
afford_city_senior=afford_city_senior.sort_values(by=['affordability'],
      ascending=False)
afford_city_both=afford_city_both.sort_values(by=['affordability'],
      ascending=False)
print("Top five most affordable cities for JUNIOR data scientists:")
print(afford_city_junior.head(5))
print("Top five most affordable cities for SENIOR data scientists:")
print(afford_city_senior.head(5))
print("Top five most affordable cities for data scientists regardless of \
level:")
print(afford_city_both.head(5))
```

Top five most affordable cities for JUNIOR data scientists:

	Country	City	State	RealCOL_Index	sal_index \
79	United Arab Emirates	Dubai	-	54.83	93.826515
109	United States	Berkeley	CA	91.48	155.100973

233	United States	Oakland	CA	90.52	151.909595
244	United States	Pittsburgh	PA	62.10	104.038924
87	United States	Ann Arbor	MI	59.82	100.209271

	affordability
79	171.122587
109	169.546320
233	167.818819
244	167.534500
87	167.518005

Top five most affordable cities for SENIOR data scientists:

	Country	City	State	RealCOL_Index	sal_index \
271	United States	San Antonio	TX	51.19	173.189575
43	India	Mumbai	-	24.82	57.152560
198	United States	Los Angeles	CA	76.98	170.358591
33	India	Bangalore	-	19.01	34.767807
302	United States	Tulsa	OK	46.68	82.265048

	affordability
271	338.326968
43	230.268169
198	221.302405
33	182.892199
302	176.231894

Top five most affordable cities for data scientists regardless of level:

	Country	City	State	RealCOL_Index	sal_index \
92	United States	San Antonio	TX	51.19	123.295635
73	United States	Los Angeles	CA	76.98	144.919438
25	India	Mumbai	-	24.82	44.533170
99	United States	Tulsa	OK	46.68	82.265048
42	United Arab Emirates	Dubai	-	54.83	93.826515

	affordability
92	240.858830
73	188.255960
25	179.424537
99	176.231894
42	171.122587

## 8 Part 7 - Comparing junior to senior affordability


If we plot the kernel density of affordability indexes to see if there are insights, an interesting insight comes out of it: Both affordability plots seem pretty normally distributed. But given that 100 is the index base for NYC, we see the following:

- For Senior people, New York is right in the middle of the distribution meaning there are as many cities less affordable than NYC than more
- However, for Junior people, the median of the distribution is shifted to the left meaning there

are more cities that are more affordable than NYC than there are those less affordable. Seems NYC hero worship of experience, and making lower level people suffer, are alive and well!

One other thing of note is that the distribution of affordability for Senior people has a long “tail” to the right, suggesting that there are outliers in the data where Senior people were well overpaid to be lured to a low cost of living city. Digging into the data, this seems to be why San Antonio found its way into the top 5 for Senior people

```
[34]: junior_scores=afford_city_junior['affordability']
junior_scores=junior_scores.rename("Junior") # rename
senior_scores=afford_city_senior['affordability']
junior_scores=junior_scores.rename("Senior") # rename
fig = sns.kdeplot(junior_scores, fill=True, color="r") # plot junior scores
fig = sns.kdeplot(senior_scores, fill=True, color="b") # plot senior scores
plt.title('City Affordability Distribution - JUNIOR Data Scientists')
plt.ylabel('Density')
plt.xlabel('Affordability Index (NYC=100)')
plt.legend(labels=["Junior DS", "Senior DS"], loc='upper right')
plt.show()
```



## 9 Part 8 - Identifying factors driving affordability

One other useful graphing analysis looks at what drives affordability in a given city for data scientists, and if it's different for Junior versus Senior people. If cost of living and salaries are in synch, there shouldn't be much to infer. But what if they are not? I looked at two possible factors:


- 1) Affordability is driven by cost-of-living - meaning, cities with lower cost of living (including housing) are more affordable even despite paying lower salaries.
- 2) Affordability is driven by salaries - meaning, cities with higher salaries are more affordable even when real cost of living is high

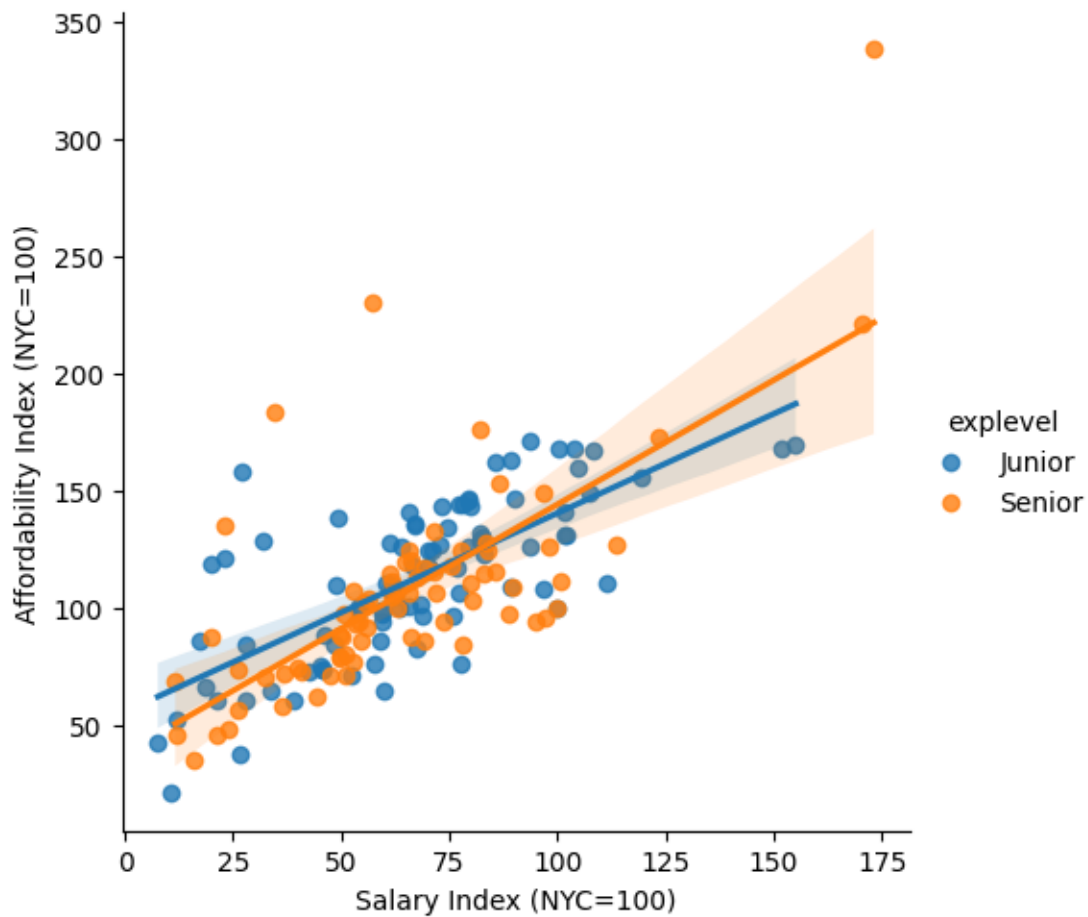
What the below plots seem to show is :

- 1) Cost of living isn't a very big driver of affordability of a place.  
The vast majority of cities are cheaper than NYC, but as many of them are less affordable (affordability index <100) as more affordable, and it doesn't seem to matter how cheap a place is to live.  
The flat trendline on the first graph shows this
- 2) On the other hand, the second plot seems to suggest that local salaries are a significant driver of affordability. The more money you make, both at a junior and senior level, the more a specific city is affordable no matter what the local cost of living is

It could be that some of my assumptions (especially assuming offshore salaries are in USD in that file) are faulty. It could also be that the data sample size is too small. But hopefully I've demonstrated the power of data analytics!

```
[35]: # Now show regression of affordability vs cost of living to see if strong
# relationship
data_to_plot=afford_city_level[['explevel', 'RealCOL_Index', 'affordability']]
sns.lmplot(x = 'RealCOL_Index',
           y = 'affordability',
           hue = 'explevel',
           data=data_to_plot)
plt.ylabel('Affordability Index (NYC=100)')
plt.xlabel('Real Cost of Living Index (NYC=100)')
plt.show()
# Finally, show regression of affordability vs salaries to see if strong
# relationship
data_to_plot=afford_city_level[['explevel', 'sal_index', 'affordability']]
sns.lmplot(x = 'sal_index',
           y = 'affordability',
           hue = 'explevel',
           data=data_to_plot)
plt.ylabel('Affordability Index (NYC=100)')
plt.xlabel('Salary Index (NYC=100)')
plt.show()
```





## 10 Thank you for listening to my TED Talk

I enjoyed this class a ton! I can't believe how much I learned in two languages I had never seen (and one, "R", that I'd never even heard of.